

Narendranath Edara

+1 (573) 466-6656 | edara.narendranath@gmail.com | [linkedin.com/in/narendranathe](https://www.linkedin.com/in/narendranathe) | narendranathe.github.io

EXPERIENCE

Data Engineer

July 2024 – Present

ExponentHR

Dallas, TX

- Built governed semantic layer on Microsoft Fabric standardizing metric definitions across BI and finance stakeholders, cutting support tickets by ~40% and query response from 12s to under 4s
- Delivered a Databricks analytics solution with infrastructure provisioned as code via Terraform, orchestrating cloud compute and storage for repeatable, automated environment delivery
- Reengineered CDC-based ETL from full-table reloads to incremental capture with merge upserts, cutting batch runtime from ~30 min to under 8 min and compute cost by ~67% by eliminating full-table I/O and data shuffling
- Engineered one-click Azure DevOps pipeline with idempotent provisioning for Contained AAG databases (restore, security, recovery, listener validation, CDC), saving ~1 hour per request across 20+ daily copy-downs
- Compressed deployment cycles from 3 months to 14 days by owning CI/CD end-to-end in Azure DevOps, restructuring cross-team handoffs that accounted for 11 weeks of idle time
- Owned payroll-critical recovery via containerized AAG failover, holding a 98% uptime SLA with sub-hour restore targets including CDC capture-job recovery after failover events

Data Engineer

Aug. 2023 – July 2024

Missouri S&T

Rolla, MO

- Engineered Azure AI Anomaly Detector pipelines selecting the optimal algorithm per time-series profile at 95%+ accuracy, catching a production memory leak 4 hours before outage
- Implemented per-service tunable thresholds filtering ~250 weekly non-actionable P3 alerts, moving signal-to-noise from 1:5 to 1:1.2 and ending alert blindness on on-call rotations
- Migrated static D-Series VMs to AKS with HPA, raising CPU utilization from 12% to 64%, consolidating 20 nodes to 4-8 dynamic, and cutting monthly Azure spend by \$3,200

Business Analyst

Mar. 2018 – Sept. 2020

Zomato

Hyderabad, India

- Built a real-time competitor analytics platform tracking pricing, delivery times, and coverage, feeding pricing changes that contributed to a 9% market-share gain in contested metros
- Optimized search relevance with ranking models using contextual signals (time, location, cuisine affinity), improving search-to-conversion across millions of daily queries
- Built an Elasticsearch enterprise search index over 100K+ documents, reducing Support Desk volume by ~80% through self-serve query resolution

EDUCATION

Missouri University of Science and Technology

Rolla, MO

Master of Science in Information Science and Technology; GPA: 4.0/4.0

Jan. 2022 – Dec. 2023

PROJECTS

JobScout - Real-Time Job Discovery Platform | *Python, Flask, SQLite, GitHub Actions, React*

2026

- Built a real-time discovery platform monitoring 130+ companies across 6 ATS providers (Greenhouse, Lever, Ashby, Workday), firing instant Discord and Telegram alerts on matched roles via ranked title and skill scoring
- Designed a \$0/month self-healing architecture on Flask, SQLite, and GitHub Actions with circuit breakers isolating failing scrapers and a one-click health-check probing DB, scrapers, and recent runs

Real-Time Portfolio Risk Analytics | *Kafka, Spark, AWS Kinesis, FastAPI, Docker*

2026

- Built a streaming risk platform ingesting market data at 47.8 TPS through Kafka with Spark 5-second windowed aggregations, computing Value-at-Risk at 95% and 99% confidence with sub-second dashboard refresh
- Containerized the end-to-end pipeline (Kafka, Spark, AWS Kinesis via boto3, FastAPI, Streamlit) on Docker Compose, exposing concentration and diversification risk metrics through a REST API

CERTIFICATIONS & PUBLICATIONS

DP-700: Microsoft Certified Data Engineer Associate | **Publication:** Sentiment Analysis for Visitor Insights – [DOI](#)

TECHNICAL SKILLS

Languages & Frameworks: Python, SQL, Bash, FastAPI, PySpark, Pytest, Git, CI/CD

Data Engineering: Spark, Kafka, Delta Lake, CDC, ETL/ELT, Data Modeling, Airflow, Data Quality, Observability

Cloud & Platforms: Azure (AKS, DevOps, Fabric), AWS (S3, Kinesis), Databricks, Terraform, Docker, Kubernetes

AI/ML: RAG, LangChain, MLflow, Feature Engineering, Model Serving, Prometheus, Grafana